

Lecture 8

Neural Networks




INE2018

Kwon, Bokyung

Neural Networks

➤ Properties?

- Uni-direction
- Convergence and divergence
- Summation
- Excitation or inhibition



< neuron synapse image: Google >

INE2018

Kwon, Bokyung

Neural Networks

➤ Networks

- Networks consists of multiple layers
- Each layer consists of a set of nodes
- Nodes are linked to other nodes in another layer


INE2018

Kwon, Bokyung

Neural Networks

➤ Structure

- A neural network model with a single hidden layer



INE2018

Kwon, Bokyung

Neural Networks

➤ Modeling

- we have:

$$Y = \beta_0 + \beta_1 H_1 + \cdots + \beta_{m-1} H_{m-1}$$

- and:

$$H_j = \alpha_{j0} + \alpha_{j1} X_1 + \cdots + \alpha_{j,p-1} X_{p-1}$$

INE2018

Kwon, Bokyung

Neural Networks

➤ Activation functions

- Identity activation function: $\sigma(z) = z$
- Logistic (sigmoid) activation function: $\sigma(z) = \begin{cases} 1 & , z \rightarrow \infty \\ 0 & , z \rightarrow -\infty \end{cases}$
- Softmax activation function: $\sigma(z) = \frac{\exp(z)}{\sum_{j=1}^n \exp(z_j)}$

INE2018

Kwon, Bokyung

Neural Networks

➤ Activation functions

- **ReLU** activation function: $\sigma(z) = \max(0, z)$
- **Leaky ReLU** activation function: $\sigma(z) = \max(\alpha z, z)$
- **Tanh** activation function: $\sigma(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$

INE2018

Kwon, Bokyung

Neural Networks

➤ Modeling

- These expressions can be generalized using an activation function.

$$Y = \sigma(\beta_0 + \sum_j^{m-1} \beta_j H_j)$$

$$H_j = \sigma(\alpha_{j0} + \sum_k^{p-1} \alpha_{jk} X_k)$$

- $\sigma(z)$: an activation function.

INE2018

Kwon, Bokyung

Neural Networks

➤ Some useful loss functions

- SSE (Sum of Squared Error)

$$SSE = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- MSE (Mean Squared Error)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

INE2018

Kwon, Bokyung

Neural Networks

➤ Some useful loss functions

- Categorical Cross Entropy

$$CCE = \sum_{x \in \chi} p(x) \log(q(x))$$

- Kullback-Leibler (KL) divergence

$$D_{KL}(P||Q) = \sum_{x \in \chi} p(x) \log\left(\frac{P(x)}{Q(x)}\right)$$

INE2018

Kwon, Bokyung


Practice

INE2018

Kwon, Bokyung

➤ Toy example

x1	x2	x3	Y
0	0	1	0
0	1	1	1
1	0	1	1
1	1	1	0



INE2018

Kwon, Bokyung

➤ Dataset

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

In [4]: x = np.array([[0,0,1],[0,1,1],[1,0,1],[1,1,1]])
y = np.array([[0],[1],[1],[0]])

print("The shape of X is ", x.shape)
print("The shape of Y is ", y.shape)

The shape of X is (4, 3)
The shape of Y is (4, 1)
```

INE2018

Kwon, Bokyung

➤ Define the activation and the loss function

```
In [5]: def act_sigmoid(z):
    return 1/(1+np.exp(-z))

def grad_sigmoid(z):
    return z*(1-z)

def sse(y, output):
    return np.sum(np.power(y-output,2))
```

INE2018

Kwon, Bokyung


➤ Working with a neural network

```
In [8]: class nn_toy:  
    def __init__(self, x, y):  
        self.x = x  
        self.y = y  
        self.alphas = np.random.rand(self.x.shape[1], 4)  
        self.betas = np.random.rand(4, 1)  
        self.output = np.zeros(self.y.shape)  
  
    def forward_process(self):  
        self.H = act_sigmoid(np.dot(self.x, self.alphas))  
        self.output = act_sigmoid(np.dot(self.H, self.betas))  
  
    def backprop_process(self):  
        grad_betas = np.dot(self.H.T, (2*(self.y-self.output)*grad_sigmoid(self.output)))  
        grad_alphas = np.dot(self.x.T, (np.dot(2*(self.y-self.output)*grad_sigmoid(self.output),  
                                              self.betas.T)*grad_sigmoid(self.H)))  
        self.alphas += grad_alphas  
        self.betas += grad_betas
```

INE2018

Kwon, Bokyung

➤ Running...



```
In [20]: my_data = nn_toy(x,y)  
  
for i in range(1500):  
    my_data.forward_process()  
    my_data.backprop_process()  
    print("SSE:", sse(y, my_data.output))
```

```
SSE: 1.3633628732530294  
SSE: 1.0620406646638603  
SSE: 0.997403335270371  
SSE: 0.9968249679842298  
SSE: 0.9966096453576148  
SSE: 0.9964027653381119  
SSE: 0.9961906927675295  
SSE: 0.9959725883569573  
SSE: 0.995748074932304  
SSE: 0.9955167780805947  
SSE: 0.995278303405698  
SSE: 0.995032259904416  
SSE: 0.9947782089482315  
SSE: 0.9945157132258386  
SSE: 0.994244310633263
```

INE2018

Kwon, Bokyung

➤ Predicted probability

```
In [21]: print("_____")
print("+      predicted      +")
print("_____")
print(my_data.output)
```

+ predicted +		Y
		0
		1
		1
		0

True values of Y

➤ Building a neural network

- we need `models.Sequential()` and `layers.Dense()`.
 - `models.Sequential()`
: helps to construct a model by stacking layers sequentially
 - `layers.Dense()`
: Dense operates element-wise activation function as:

$$\sigma\left(\sum_{i=1}^K \beta_i \cdot x_i + b_l\right)$$

➤ Building a neural network

- `..fit(x, y, epochs, verbose)`
 - x, y: x and y in arrays
 - epochs: training epochs
 - verbose =0 or 1 or 2 : 0 is silent, 1 shows progress bar,
2 as a single line per each epoch
- `.compile(loss, optimizer)`
 - loss: loss functions such as ‘mse’, ‘categorical_crossentropy’
 - optimizer: training algorithm such as ‘sgd’, ‘adam’

➤ Hand-written digits Dataset (Kaynak, C. 1995)

- we can use “digits” dataset for classification.
- we can use a dataset from *keras*

```
In [2]: from keras.datasets import mnist
train_set, test_set = mnist.load_data()
x_train, y_train = train_set
x_test, y_test = test_set

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 1s 0us/step
```

INE2018

Kwon, Bokyung

➤ Observation by the label

```
In [4]: fig, ((ax1, ax2, ax3, ax4, ax5), (ax6, ax7, ax8, ax9, ax10)) = plt.subplots(2,5, figsize=(10,5))

for idx, ax in enumerate([ax1, ax2, ax3, ax4, ax5, ax6, ax7, ax8, ax9, ax10]):
    for i in range(1000):
        if y_test[i]== idx:
            ax.imshow(x_test[i],cmap='gray')
            ax.grid(False)
            ax.set_xticks([])
            ax.set_yticks([])
            break

plt.tight_layout()
plt.show()
```

INE2018

Kwon, Bokyung

➤ Observation



INE2018

Kwon, Bokyung

➤ Data reshaping and normalization

```
In [9]: print(x_train.shape, x_test.shape)
```

```
(60000, 28, 28) (10000, 28, 28)
```

```
In [10]: x_train_vec = x_train.reshape((x_train.shape[0], x_train.shape[1]*x_train.shape[2]))
```

```
x_test_vec = x_test.reshape((x_test.shape[0], x_test.shape[1]*x_test.shape[2]))
```

```
print("The dim of x_train: ", x_train_vec.shape)
```

```
print("The dim of x_test: ", x_test_vec.shape)
```

```
The dim of x_train: (60000, 784)
```

```
The dim of x_test: (10000, 784)
```

```
In [6]: x_train_vec = x_train_vec/255.
```

```
x_test_vec = x_test_vec/255.
```

INE2018

Kwon, Bokyung

➤ Label?

```
In [3]: y_train  
Out[3]: array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)  
  
In [12]: from keras.utils.np_utils import to_categorical  
  
y_train = to_categorical(y_train, 10)  
y_test = to_categorical(y_test, 10)  
  
In [13]: y_train  
Out[13]: array([[0., 0., 0., ..., 0., 0., 0.],  
                [1., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                ...,  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 1., 0.]], dtype=float32)
```

INE2018

Kwon, Bokyung

➤ Constructing a NN using API (1)

```
In [4]: import tensorflow as tf  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense  
  
num_classes=10  
my_model = Sequential()  
my_model.add(Dense(units = 20, input_shape =(x_train_vec.shape[1]), activation='relu'))  
my_model.add(Dense(units = num_classes, activation='softmax'))
```

INE2018

Kwon, Bokyung

➤ Summary()

```
In [5]: my_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 20)	15700
dense_1 (Dense)	(None, 10)	210
<hr/>		
Total params: 15910 (62.15 KB)		
Trainable params: 15910 (62.15 KB)		
Non-trainable params: 0 (0.00 Byte)		

INE2018

Kwon, Bokyung

➤ Constructing a NN using API (2)

```
In [20]: from keras.models import Sequential  
from keras.layers import Dense, Dropout, Activation, Flatten  
from keras.optimizers import SGD, Adam, RMSprop
```

```
In [22]: num_classes=10  
my_model2 = Sequential()  
my_model2.add(Dense(units = 512, input_shape =(x_train_vec.shape[1],)))  
my_model2.add(Activation('relu'))  
my_model2.add(Dropout(0.2))  
my_model2.add(Dense(units = 512))  
my_model2.add(Activation('relu'))  
my_model2.add(Dropout(0.2))  
my_model2.add(Dense(units = num_classes))  
my_model2.add(Activation('softmax'))
```

INE2018

Kwon, Bokyung

```
In [23]: my_model2.summary()
Model: "sequential_2"
-----
```

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 512)	401920
activation_3 (Activation)	(None, 512)	0
dropout_2 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 512)	262656
activation_4 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 10)	5130
activation_5 (Activation)	(None, 10)	0

```
=====
Total params: 669706 (2.55 MB)
Trainable params: 669706 (2.55 MB)
Non-trainable params: 0 (0.00 Byte)
-----
```

INE2018

Kwon, Bokyung

➤ Compiling and fitting (my_model 2)

```
In [24]: my_model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
my_model2.fit(x_train_vec, y_train, epochs=10, verbose=1)

Epoch 1/10
1875/1875 [=====] - 23s 12ms/step - loss: 0.2120 - accuracy: 0.9354
Epoch 2/10
1875/1875 [=====] - 21s 11ms/step - loss: 0.1058 - accuracy: 0.9676
Epoch 3/10
1875/1875 [=====] - 22s 12ms/step - loss: 0.0788 - accuracy: 0.9751
Epoch 4/10
1875/1875 [=====] - 23s 12ms/step - loss: 0.0635 - accuracy: 0.9798
Epoch 5/10
1875/1875 [=====] - 21s 11ms/step - loss: 0.0568 - accuracy: 0.9826
Epoch 6/10
1875/1875 [=====] - 21s 11ms/step - loss: 0.0495 - accuracy: 0.9843
Epoch 7/10
1875/1875 [=====] - 21s 11ms/step - loss: 0.0467 - accuracy: 0.9855
Epoch 8/10
1875/1875 [=====] - 21s 11ms/step - loss: 0.0403 - accuracy: 0.9873
Epoch 9/10
1875/1875 [=====] - 21s 11ms/step - loss: 0.0391 - accuracy: 0.9880
Epoch 10/10
1875/1875 [=====] - 21s 11ms/step - loss: 0.0384 - accuracy: 0.9879
```

```
Out[24]: <keras.callbacks.History at 0x1d9617f7f90>
```

INE2018

Kwon, Bokyung

➤ Prediction

```
In [25]: output = my_model2.predict(x_test_vec)
313/313 [=====] - 1s 3ms/step

In [26]: output[0,:]

Out[26]: array([4.5980282e-14, 7.1095255e-11, 2.2075770e-10, 7.0837801e-11,
       2.3529572e-11, 1.1822084e-13, 1.8544796e-21, 1.0000000e+00,
       1.8453263e-14, 8.8900443e-09], dtype=float32)
```